


# **Machine Learning Workshop:**

## Deep Convolutional Neural Networks

# Why Deep Learning?

- Deep learning enables machine learning to not only learn how features can be used for classification, but also learn the features themselves from the most basic representations
- Deep learning is a process of learning simple patterns in the data and building off of those patterns to learn more complex ones

# Deep Learning for Facial Recognition



**Input Layer:**

Different Intensities of Pixels

**Layer 1:** Learn simple shapes

**Layer 2:** Learn More Complex Shapes and objects

**Layer 3:** Learn compositions of complex shapes that define a face

(Andrew Ng)

# Global vs Local Feature Learning

A typical classifier would look at each word in this sentence independently and try to identify a relationship between the frequency of words with the class label associated with it

Is identical to:

it A relationship look classifier at associated each in this try to a between sentence and the would frequency of with typical the class label with words it independently identify word

For most encodings and machine learning classifiers.  
(e.g., word frequency)

# Global vs Local Feature Learning

A typical classifier would look at each word in this sentence independently and try to identify a relationship between the frequency of words with the class label associated with it

<b>a</b>	<b>it</b>	<b>try</b>	<b>identify</b>	<b>...</b>
2	1	1	1	...

it A relationship look classifier at associated each in this try to a between sentence and the would frequency of with typical the class label with words it independently identify word

<b>a</b>	<b>it</b>	<b>try</b>	<b>identify</b>	<b>...</b>
2	1	1	1	...

# Global vs Local Feature Learning

But when we read, we look at neighbors to understand the context and meaning:

A typical classifier would look at each word in this sentence independently and try to identify a relationship between the frequency of words with the class label associated with it

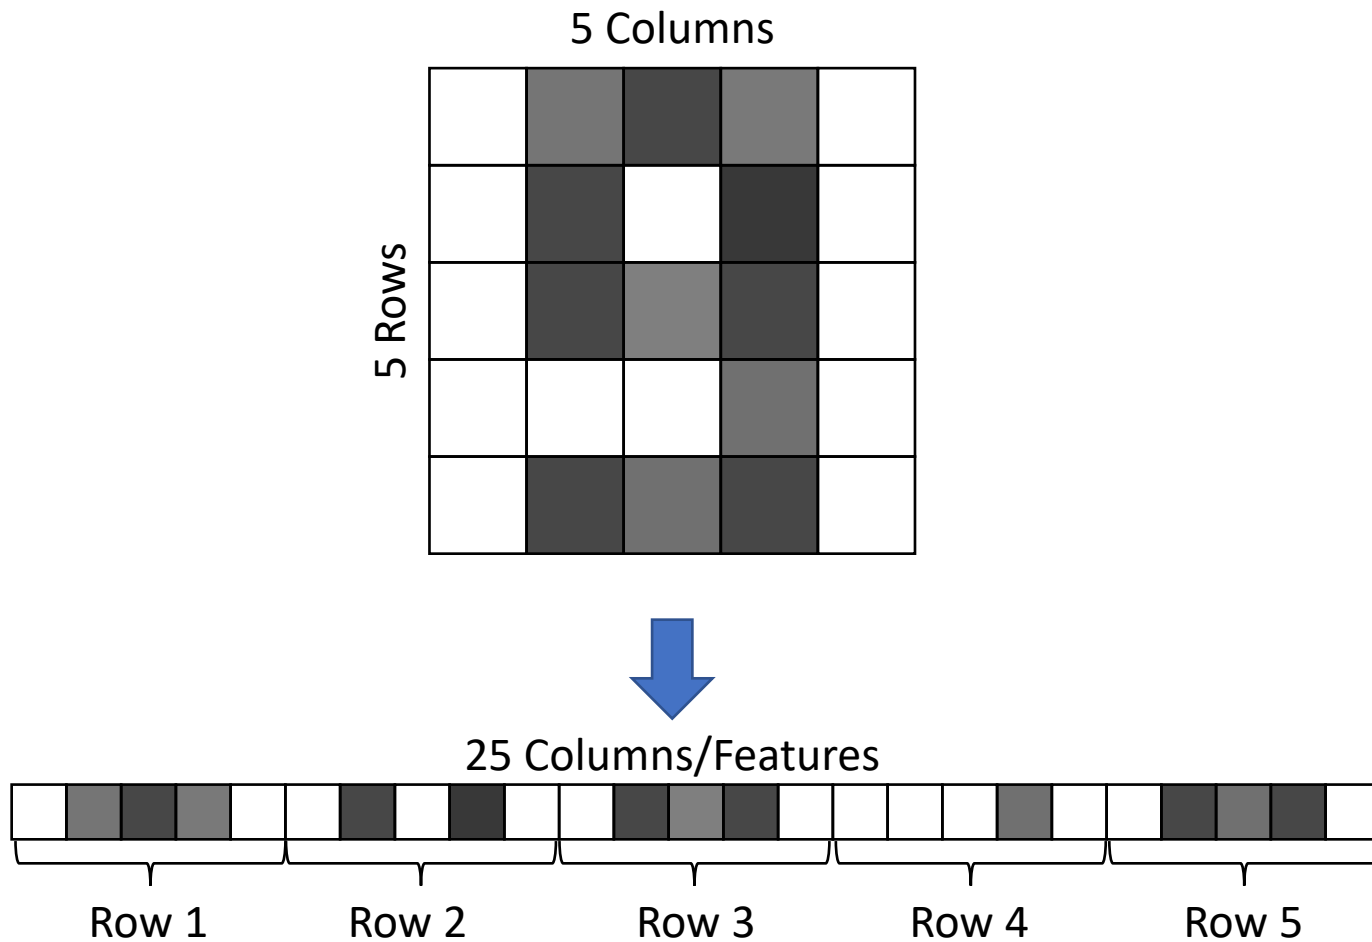
The position of the words with respect to one another is just as important as the words themselves!

**Classifiers that can learn local contexts:**

Deep Convolutional Neural Networks

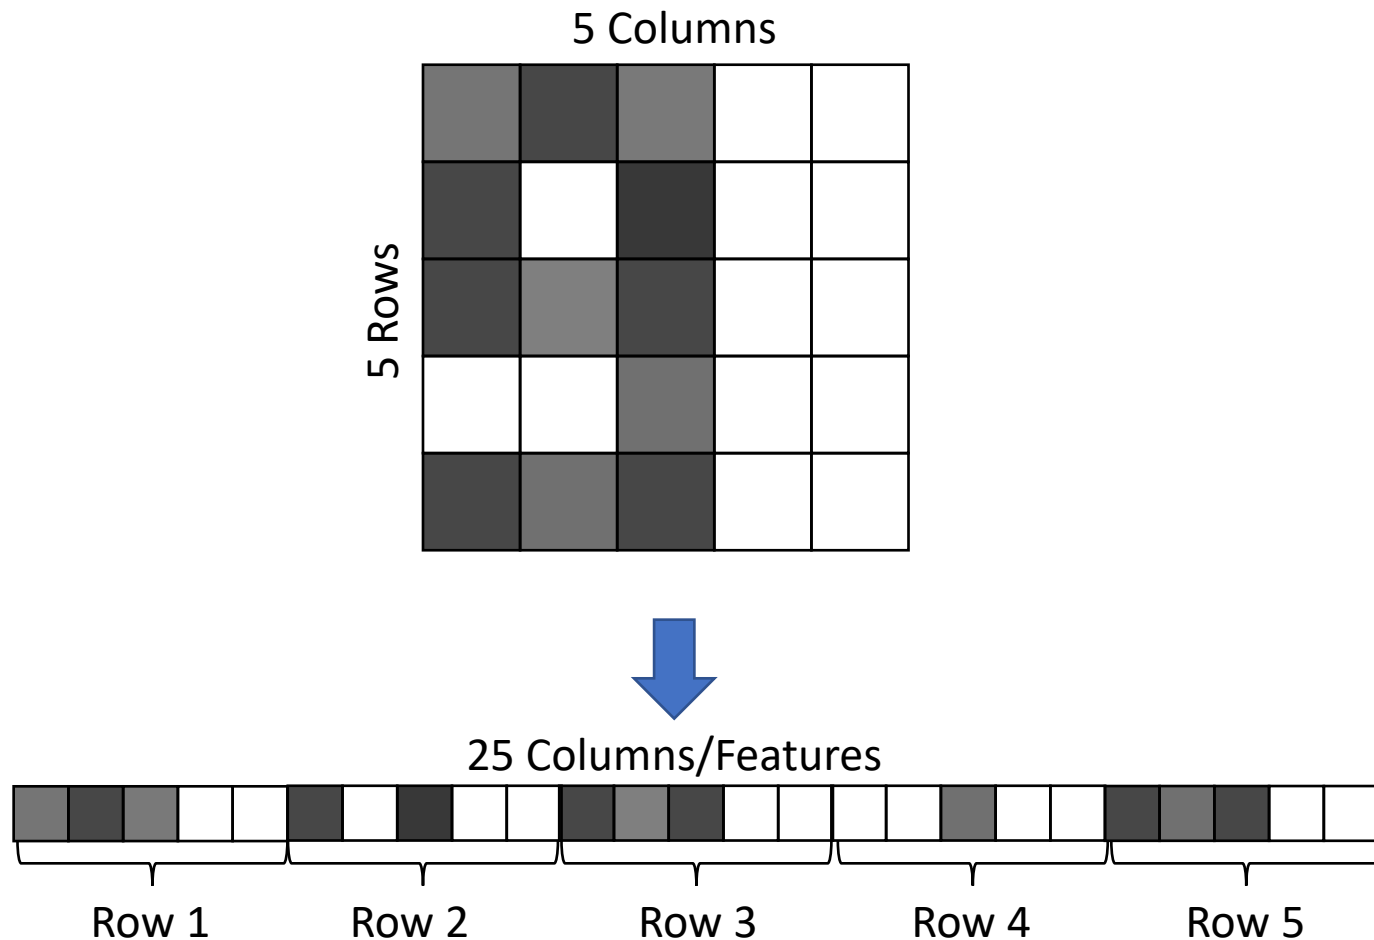
Recurrent Neural Networks

# Revisiting our Number Encoding:




This is known as **“Flattening”**

**Problem:** Each number can be shifted or represented slightly differently, increasing both model complexity and training data to capture all variations





**Problem:** Each number can be shifted or represented slightly differently, increasing both model complexity and training data to capture all variations



We can start to deconstruct our number into components, and then use these components to predict the number being represented by the image!

# The Convolution Function

The convolution function transforms the data by learning patterns from local neighborhoods in the data and generating a new representation of the data using patterns observed


**Kernel dimensions:** The size of the window/tensor used to examine data points in close proximity

**# Filters:** The number of filters to train on the data


**Strides:** How the window will translate/slide across the data to generate the new output

# Convolution Example

Input:




Output:



# Convolution Example

Input:

Kernel  
Window/  
Tensor




$$\begin{pmatrix} 0.0 & 1.0 & 0.0 \\ 0.0 & 0.5 & 1.0 \\ 0.5 & 0.3 & 0.5 \end{pmatrix} * \begin{pmatrix} w1 & w2 & w3 \\ w4 & w5 & w6 \\ w7 & w8 & w9 \end{pmatrix}$$

Observed

Filter 1

$$\begin{aligned} &0.0 * w1 + 1.0 * w2 + 0.0 * w3 + \\ &0.0 * w4 + 0.5 * w5 + 1.0 * w6 + \\ &0.5 * w7 + 0.3 * w8 + 0.5 * w9 \end{aligned}$$


Output:



Number of Filters (n=3)

# Convolution Example

Input:




$$\begin{pmatrix} 0.0 & 1.0 & 0.0 \\ 0.0 & 0.5 & 1.0 \\ 0.5 & 0.3 & 0.5 \end{pmatrix} * \begin{pmatrix} w1 & w2 & w3 \\ w4 & w5 & w6 \\ w7 & w8 & w9 \end{pmatrix}$$

Observed

Filter 1

$$\begin{aligned} &0.0 * w1 + 1.0 * w2 + 0.0 * w3 + \\ &0.0 * w4 + 0.5 * w5 + 1.0 * w6 + \\ &0.5 * w7 + 0.3 * w8 + 0.5 * w9 \end{aligned}$$


Output:



Number of Filters (n=3)

# Convolution Example

Input:



$$\begin{pmatrix} 0.0 & 1.0 & 0.0 \\ 0.0 & 0.5 & 1.0 \\ 0.5 & 0.3 & 0.5 \end{pmatrix} * \begin{pmatrix} w1 & w2 & w3 \\ w4 & w5 & w6 \\ w7 & w8 & w9 \end{pmatrix}$$


Observed

Filter 2



$$\begin{aligned} &0.0 * w1 + 1.0 * w2 + 0.0 * w3 + \\ &0.0 * w4 + 0.5 * w5 + 1.0 * w6 + \\ &0.5 * w7 + 0.3 * w8 + 0.5 * w9 \end{aligned}$$


Output:



Number of Filters (n=3)

# Convolution Example

Input:



$$\begin{pmatrix} 0.0 & 1.0 & 0.0 \\ 0.0 & 0.5 & 1.0 \\ 0.5 & 0.3 & 0.5 \end{pmatrix} * \begin{pmatrix} w1 & w2 & w3 \\ w4 & w5 & w6 \\ w7 & w8 & w9 \end{pmatrix}$$


Observed

Filter 3



$$\begin{aligned} &0.0 * w1 + 1.0 * w2 + 0.0 * w3 + \\ &0.0 * w4 + 0.5 * w5 + 1.0 * w6 + \\ &0.5 * w7 + 0.3 * w8 + 0.5 * w9 \end{aligned}$$


Output:




Number of Filters (n=3)

# Convolution Example

Input:




Output:






# Convolution Example

Input:




Output:




# Convolution Example

Input:




Output:




# Convolution Example

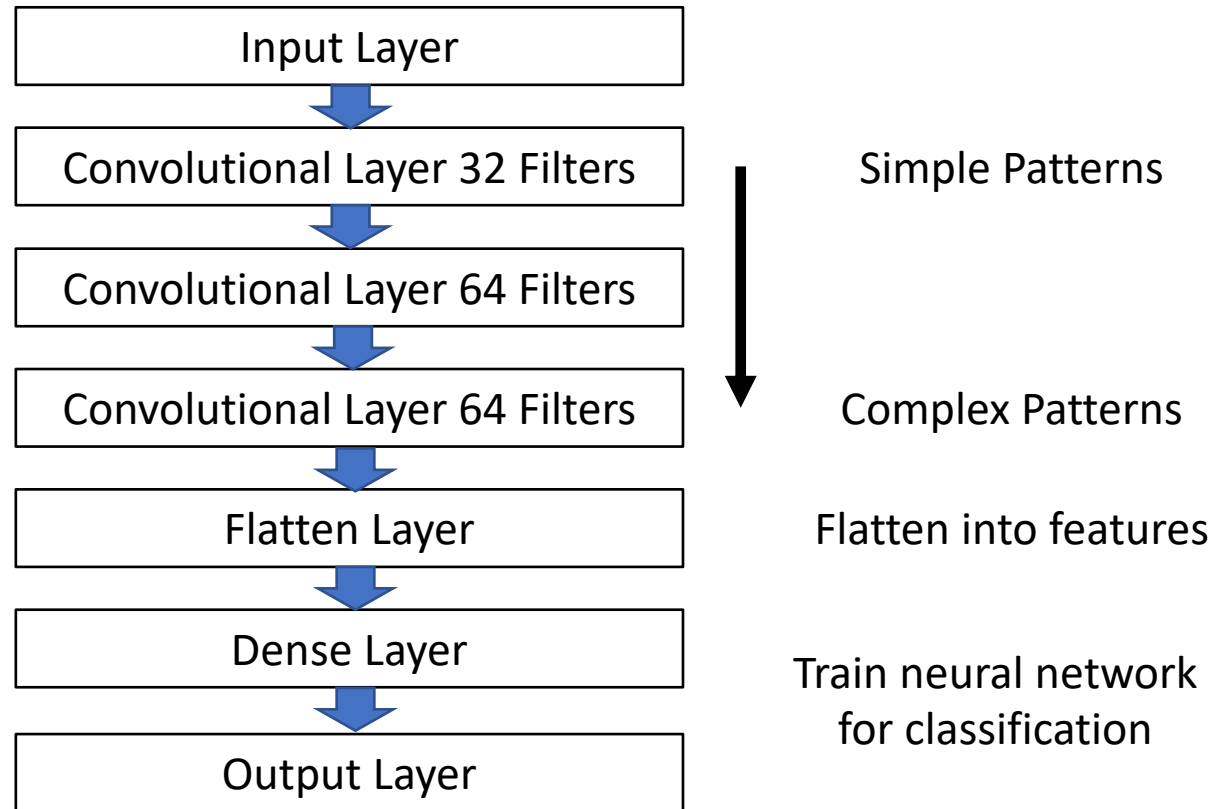
Input:



Output:

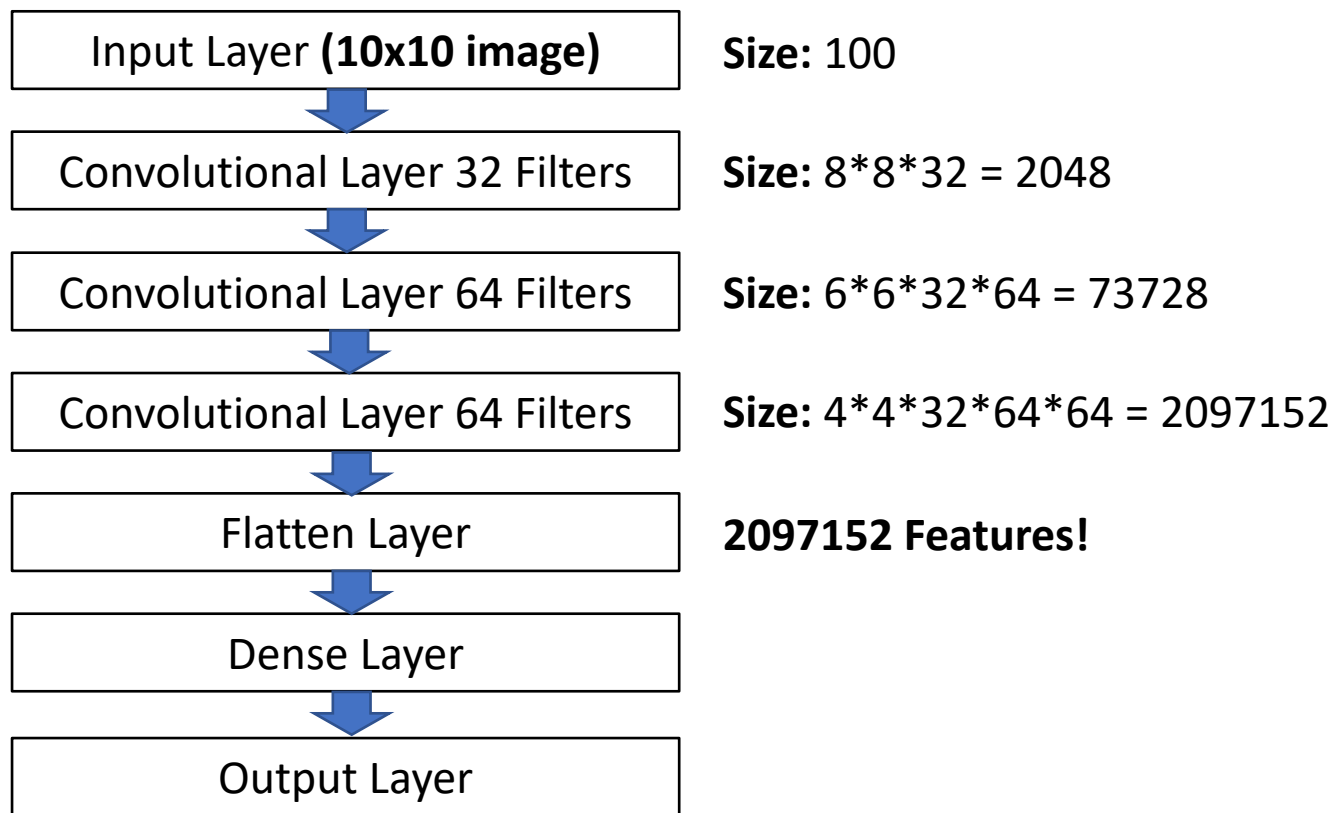


# A Typical Deep CNN Architecture




# Problem with convolution: The number of features grows as we go deeper in the network

Assuming 3x3  
kernels for each  
convolutional  
layer

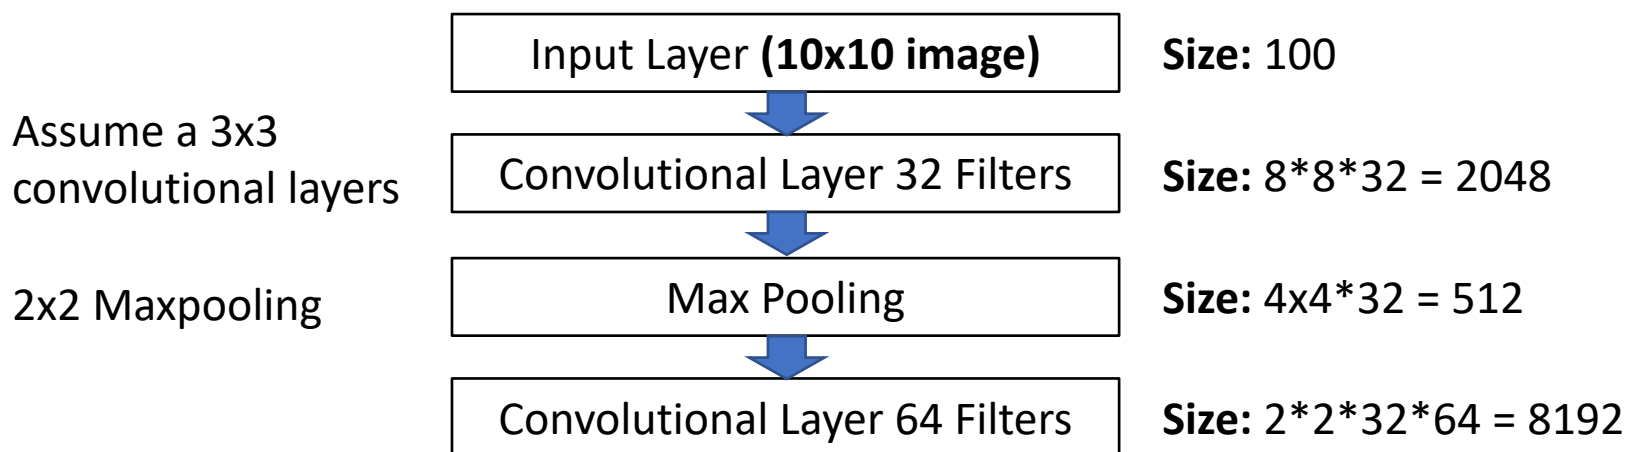


# Max Pooling

**Basic Idea:** Reduce the the dimensions of the data by taking the maximum value within a specified window



# Problem with convolution: The number of features grows as we go deeper in the network



We actually reach our limit on convolutions sooner and are forced to stop here and move to flattening into features and training a dense neural network

# Overcoming Overfitting with Dropout


**Basic idea:** randomly set a certain percentage of features to 0, forcing the machine learning algorithm to learn consistent patterns in the data

“Soon after that I went to my bank. The tellers kept changing and I asked one of them why. He said he didn't know but they got moved around a lot. I figured it must be because it would require cooperation between employees to successfully defraud the bank. This made me realize that randomly removing a different subset of neurons on each example would prevent conspiracies and thus reduce overfitting.”

– Geoff Hinton



# Dropout Layer Example



Drop out layers add random noise to the data, making it harder to overfit

# Keras Exercise

# Three Key Questions to Ask for Machine Learning

1. What is the input data? What is the ground truth?
2. What are the features/data? How is the data encoded?
3. What model will be used? What will be the architecture?